

```

# Import modules
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression

import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# Load dataset
digits = datasets.load_digits()

# Create data and label arrays
X = digits.data
y = digits.target

# Split into training and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.4)

df = pd.DataFrame(columns = 'train test'.split())
for i in range(1,8):
    kNN = KNeighborsClassifier(n_neighbors=i)
    kNN.fit(X_train, y_train)
    tr = kNN.score(X_train,y_train)
    te = kNN.score(X_test,y_test)
    df.loc[i] = [tr, te]
df

```

/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the

default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
In SciPy 1.11.0, this behavior will change: the default value of
`keepdims` will become False, the `axis` over which the statistic is
taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
In SciPy 1.11.0, this behavior will change: the default value of
`keepdims` will become False, the `axis` over which the statistic is
taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
In SciPy 1.11.0, this behavior will change: the default value of
`keepdims` will become False, the `axis` over which the statistic is
taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
In SciPy 1.11.0, this behavior will change: the default value of
`keepdims` will become False, the `axis` over which the statistic is
taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
In SciPy 1.11.0, this behavior will change: the default value of
`keepdims` will become False, the `axis` over which the statistic is
taken will be eliminated, and the value None will no longer be
accepted. Set `keepdims` to True or False to avoid this warning.
```

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py:228: FutureWarning:
Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along.
```

In SciPy 1.11.0, this behavior will change: the default value of `'keepdims'` will become `False`, the `'axis'` over which the statistic is taken will be eliminated, and the value `None` will no longer be accepted. Set `'keepdims'` to `True` or `False` to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `'skew'`, `'kurtosis'`), the default behavior of `'mode'` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `'keepdims'` will become `False`, the `'axis'` over which the statistic is taken will be eliminated, and the value `None` will no longer be accepted. Set `'keepdims'` to `True` or `False` to avoid this warning.

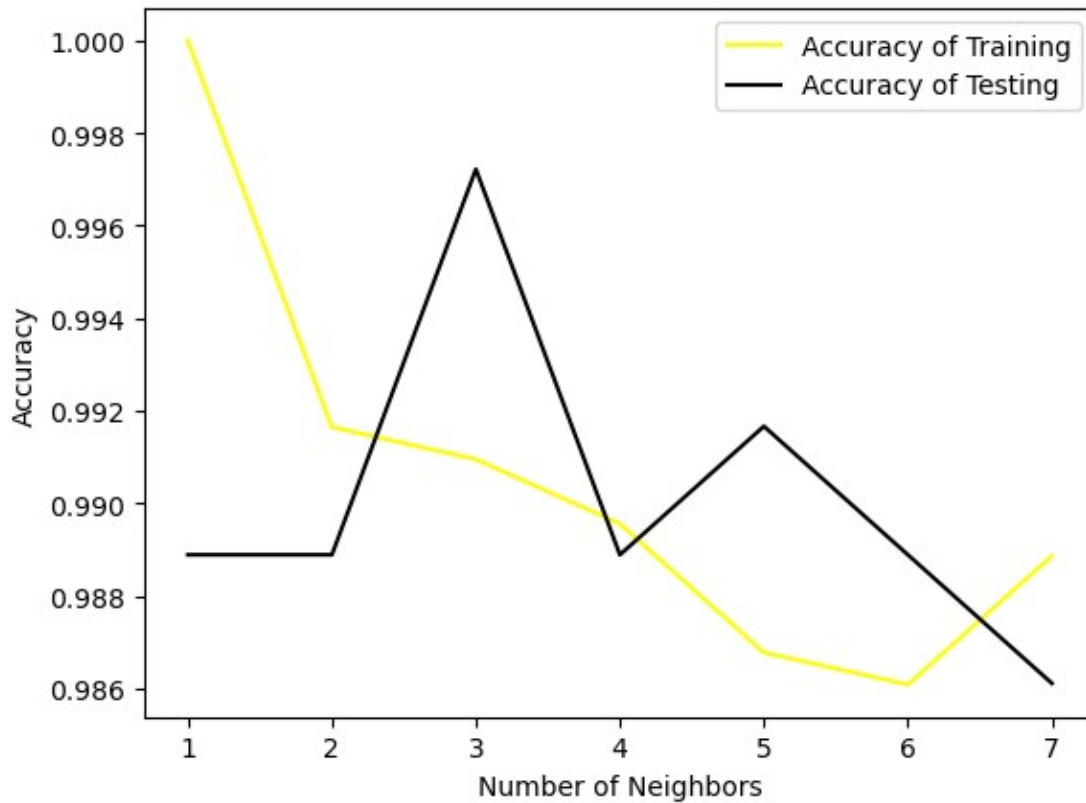
```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

	train	test
1	1.000000	0.988889
2	0.991649	0.988889
3	0.990953	0.997222
4	0.989562	0.988889
5	0.986778	0.991667
6	0.986082	0.988889
7	0.988866	0.986111

*#plot graph of train and test*

```
plt.plot(df['train'],color = 'yellow', label = 'Accuracy of Training')
plt.plot(df['test'],color = 'black', label = 'Accuracy of Testing')
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.legend()
```

<matplotlib.legend.Legend at 0x7ff2db64b7f0>



```
svm = SVC(kernel='linear')
svm.fit(X_train,y_train)
y_pred = svm.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	75
1	0.97	0.97	0.97	78
2	0.99	1.00	0.99	71
3	1.00	0.99	0.99	70
4	0.93	1.00	0.96	65
5	0.99	0.97	0.98	71
6	1.00	0.97	0.99	79
7	1.00	0.98	0.99	82
8	0.97	0.94	0.95	63
9	0.93	0.97	0.95	65
accuracy			0.98	719
macro avg	0.98	0.98	0.98	719
weighted avg	0.98	0.98	0.98	719

```
svm = SVC(kernel='poly')
svm.fit(X_train,y_train)
```

```
y_pred = svm.predict(X_test)
print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
0	1.00	0.99	0.99	75
1	0.99	1.00	0.99	78
2	0.99	1.00	0.99	71
3	1.00	0.99	0.99	70
4	0.98	1.00	0.99	65
5	0.99	0.97	0.98	71
6	1.00	0.99	0.99	79
7	1.00	0.99	0.99	82
8	0.98	1.00	0.99	63
9	0.95	0.97	0.96	65
accuracy			0.99	719
macro avg	0.99	0.99	0.99	719
weighted avg	0.99	0.99	0.99	719

```
svm = SVC(kernel='rbf')
svm.fit(X_train,y_train)
y_pred = svm.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	75
1	0.96	1.00	0.98	78
2	0.99	1.00	0.99	71
3	1.00	0.99	0.99	70
4	0.98	1.00	0.99	65
5	0.99	0.97	0.98	71
6	1.00	1.00	1.00	79
7	0.99	0.99	0.99	82
8	1.00	0.97	0.98	63
9	0.95	0.95	0.95	65
accuracy			0.99	719
macro avg	0.99	0.99	0.99	719
weighted avg	0.99	0.99	0.99	719

```
bayes = GaussianNB()
bayes.fit(X_train, y_train)
y_pred = bayes.predict(X_test)
conf_matrix = confusion_matrix(y_test,y_pred)
sns.heatmap(conf_matrix, annot= True)
```

<AxesSubplot:>



```
kNN = KNeighborsClassifier(n_neighbors=3)
kNN.fit(X_train, y_train)
a = kNN.score(X_train,y_train)
```

```
svm = SVC(kernel='poly')
svm.fit(X_train,y_train)
b = svm.score(X_train,y_train)
```

```
bayes = GaussianNB()
bayes.fit(X_train, y_train)
c = bayes.score(X_train,y_train)
```

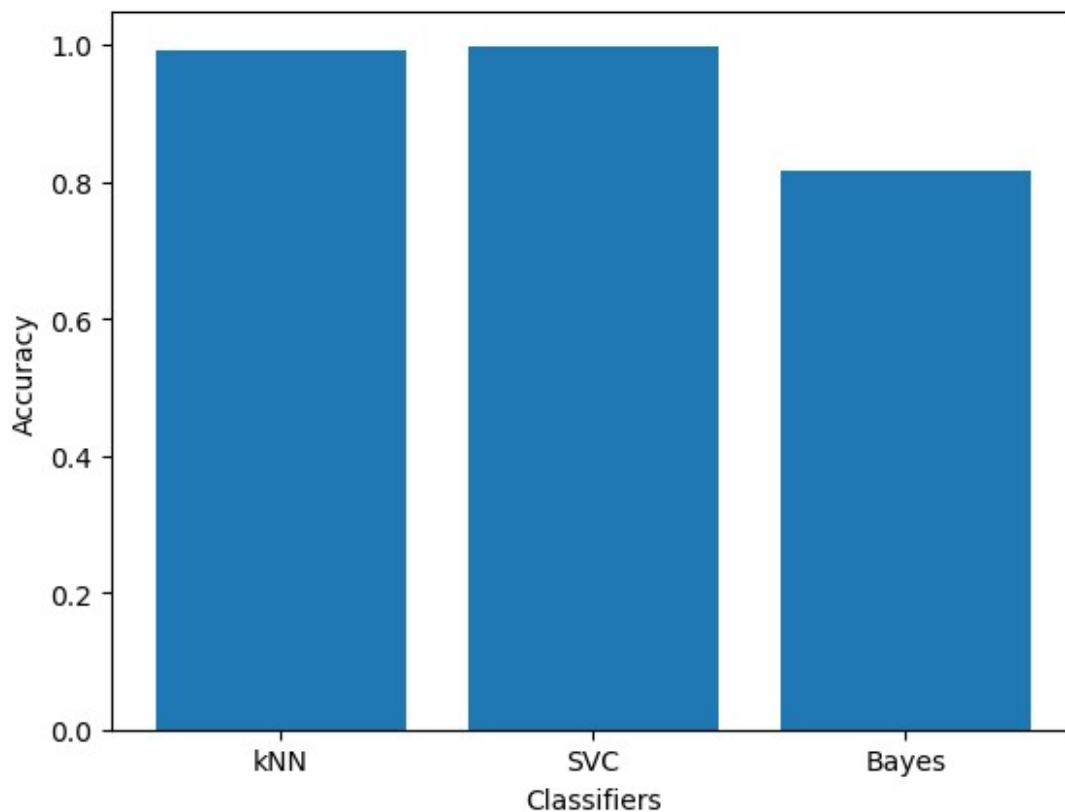
```
d = np.array(["kNN", "SVC", "Bayes"])
e = np.array([a,b,c])
```

```
plt.bar(d,e)
plt.xlabel("Classifiers")
plt.ylabel("Accuracy")
plt.show()
```

/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is

taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```



```
df = pd.read_csv("/Users/saikatduttatanu/Desktop/TRIMESTER 1,
2023/Data Analytics/gapminder.csv")
df.head()
```

	population	fertility	HIV	C02	BMI_male	GDP	BMI_female
life \							
0	34811059	2.73	0.1	3.328945	24.59620	12314	129.9049
75.3							
1	19842251	6.43	2.0	1.474353	22.25083	7103	130.1247
58.3							
2	40381860	2.24	0.5	4.785170	27.50170	14646	118.8915
75.5							
3	2975029	1.40	0.1	1.804106	25.35542	7383	132.8108
72.5							
4	21370348	1.96	0.1	18.016313	27.56373	41312	117.3755
81.5							


	child_mortality	Region
0	29.5	Middle East & North Africa
1	192.0	Sub-Saharan Africa
2	15.4	America



```
3          20.0      Europe & Central Asia
4          5.2      East Asia & Pacific
```

```
sns.heatmap(df.corr(), cmap = "RdYlGn", annot=True)
```

```
<AxesSubplot:>
```



```
# Linear Regression
```

```
x = df["fertility"]
```

```
y = df["life"]
```

```
regression = LinearRegression()
```

```
regression.fit(x[:, None], y)
```

```
print(regression.score(x[:, None], y))
```

```
fig, ax = plt.plot(x, y, 'bo', x, regression.predict(x[:, None]), '-k')
```


```
0.6192442167740035
```

```
/var/folders/sl/4pwtl8ws3psb02vg3snwbjd40000gn/T/
```

```
ipykernel_89987/4038681955.py:6: FutureWarning: Support for multi-dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be removed in a future version. Convert to a numpy array before indexing
```

instead.

```
regression.fit(x[:, None], y)
/var/folders/sl/4pwtl8ws3psb02vg3snwbjd40000gn/T/ipykernel_89987/40386
81955.py:7: FutureWarning: Support for multi-dimensional indexing
(e.g. `obj[:, None]`) is deprecated and will be removed in a future
version. Convert to a numpy array before indexing instead.
print(regression.score(x[:, None], y))
/var/folders/sl/4pwtl8ws3psb02vg3snwbjd40000gn/T/ipykernel_89987/40386
81955.py:9: FutureWarning: Support for multi-dimensional indexing
(e.g. `obj[:, None]`) is deprecated and will be removed in a future
version. Convert to a numpy array before indexing instead.
fig, ax = plt.plot(x, y, 'bo', x, regression.predict(x[:, None]), '-
k')
```



*#5 fold cross-validation*

```
cross_val_score(regression, x[:,None], y, cv=5)
```

```
/var/folders/sl/4pwtl8ws3psb02vg3snwbjd40000gn/T/
ipykernel_89987/3833724394.py:2: FutureWarning: Support for multi-
dimensional indexing (e.g. `obj[:, None]`) is deprecated and will be
removed in a future version. Convert to a numpy array before indexing
instead.
```

```
cross_val_score(regression, x[:,None], y, cv=5)
```

```
array([0.71001079, 0.75007717, 0.55271526, 0.547501 , 0.52410561])
```

```


# Linear Regression with all features
x =
df[["fertility", "HIV", "CO2", "BMI_male", "GDP", "BMI_female", "child_mortality"]]
y = df["life"]

regression1 = LinearRegression()
regression1.fit(x, y)
print(regression1.score(x, y))

plt.plot(x, y, 'bo', x, regression1.predict(x), '-k')
plt.show()

0.8974995317332959

```



```

#5 fold cross-validation

cross_val_score(regression1, x, y, cv=5)

array([0.81621881, 0.8322471 , 0.90868335, 0.81325568, 0.94404223])

#Compare the above linear regressions.
# II 2. Credit card
df1 = pd.read_csv("/Users/saikatduttatanu/Desktop/TRIMESTER 1, 2023/Data Analytics/creditcard.csv")
print(df1.shape)
df1.head()

```

(284807, 31)

	Time	V1	V2	V3	V4	V5	V6
V7 \							
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
0.239599							
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
0.078803							
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
0.791461							
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
0.237609							
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921
0.592941							

	V8	V9	...	V21	V22	V23	V24
V25 \							
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
0.128539							
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
0.167170							
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
0.327642							
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
0.647376							
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267
0.206010							

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

*# Logistic Regression*

X = df1.drop("Class",axis = 1)

y = df1.Class

log\_regression = LogisticRegression()

log\_regression.fit(X, y)

y\_pred = log\_regression.predict(X)

conf\_matrix = confusion\_matrix(y, y\_pred)

sns.heatmap(conf\_matrix, annot=True)

/Users/saikatduttatanu/Desktop/ANACONDA/anaconda3/lib/python3.9/site-packages/sklearn/linear\_model/\_logistic.py:814: ConvergenceWarning:

```
lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

<AxesSubplot:>

